

# Diffusion Model

Machine Learning pour Physicien - Projet

---

Clément, Grégoire, Nathan


January 16, 2025

# Modèle de diffusion

---

# Présentation générale

Les modèles de diffusion permettent de générer des images distribuées **de la même manière** qu'un dataset donnée.



Dataset



Images générées

**Figure 1:** Source : Denoising Diffusion Probabilistic Models, Ho. and Al.

# Application en physique

---


Eliot si tu peux rajouter quelques trucs ici, genre des exemples avec des images

# **Denoising Diffusion Probabilistic Models, théorie**

---

# General Idea

Consider the set of **hand-written digits**  $D$ . Can you give a probability distribution  $q$  such that  $x \sim q(x)$  ?




**Figure 2:** Source: ludwig.ai

# General Idea

---


Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**. Consider the following process:



# General Idea

---


Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**. Consider the following process:






# General Idea

Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**. Consider the following process:




# General Idea

Consider the set of hand-written digits  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a clever way to sample hand-written digits. Consider the following process:




# General Idea

Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**. Consider the following process:




# General Idea

Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**. Consider the following process:



# General Idea


Consider the set of **hand-written digits**  $D$ . It is hard to find  $q$  such that  $x \sim q(x)$ , we need a **clever way to sample hand-written digits**. Consider the following process:



Formally:  $q(x_{t+1} | x_t) := \mathcal{N}(x_{t+1}; \sqrt{1 - \beta_t}x_t, \beta_t I)$  for some schedule  $(\beta_t)_t$ . Can we **learn to reverse this process** ?


# What we want to learn

Given a noisy image  $x_t$ , we train a model to predict  $x_{t-1}$ .



# What we want to learn


Given a noisy image  $x_t$ , we train a model to predict  $x_{t-1}$ .



- Given a data image  $x_0$ , we sample  $(x_t)_{1:T}$  according to  $q(x_{1:T} | x_0) := \prod_{t=1}^T q(x_t | x_{t-1})$ ,

# What we want to learn

Given a noisy image  $x_t$ , we **train a model to predict**  $x_{t-1}$ .



- Given a **data image**  $x_0$ , we sample  $(x_t)_{1:T}$  according to  $q(x_{1:T} | x_0) := \prod_{t=1}^T q(x_t | x_{t-1})$ ,
- Given a **noisy image**  $x_t$  and  $t$ , we sample according to  $p_\theta(x_{t-1} | x_t) := \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$ .



# Decreasing training data generation cost

---

Given a data image  $x_0$ , compute  $x_t$  takes  $t$  sampling on  $q$ . But a simple trick, allows to do only one.

# Decreasing training data generation cost

Given a data image  $x_0$ , compute  $x_t$  takes  $t$  sampling on  $q$ . But a simple trick, allows to do only one.

Remember that  $q(x_{t+1} | x_t) := \mathcal{N}(x_{t+1}; \sqrt{1 - \beta_t}x_t, \beta_t I)$ . Let  $\alpha_t = 1 - \beta_t$  and  $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$ .

# Decreasing training data generation cost

Given a data image  $x_0$ , compute  $x_t$  takes  $t$  sampling on  $q$ . But a simple trick, allows to do only one.

Remember that  $q(x_{t+1} | x_t) := \mathcal{N}(x_{t+1}; \sqrt{1 - \beta_t}x_t, \beta_t I)$ . Let  $\alpha_t = 1 - \beta_t$  and  $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$ .

$$x_t = \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon_{t-1}$$

# Decreasing training data generation cost

Given a data image  $x_0$ , compute  $x_t$  takes  $t$  sampling on  $q$ . But a simple trick, allows to do only one.

Remember that  $q(x_{t+1} | x_t) := \mathcal{N}(x_{t+1}; \sqrt{1 - \beta_t}x_t, \beta_t I)$ . Let  $\alpha_t = 1 - \beta_t$  and  $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$ .

$$\begin{aligned}x_t &= \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon_{t-1} \\&= \sqrt{\alpha_t}\sqrt{\alpha_{t-1}}x_{t-2} + \sqrt{\alpha_t}\sqrt{1 - \alpha_{t-1}}\epsilon_{t-2} + \sqrt{1 - \alpha_t}\epsilon_{t-1}\end{aligned}$$

# Decreasing training data generation cost

Given a data image  $x_0$ , compute  $x_t$  takes  $t$  sampling on  $q$ . But a simple trick, allows to do only one.

Remember that  $q(x_{t+1} | x_t) := \mathcal{N}(x_{t+1}; \sqrt{1 - \beta_t}x_t, \beta_t I)$ . Let  $\alpha_t = 1 - \beta_t$  and  $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$ .

$$\begin{aligned}x_t &= \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon_{t-1} \\&= \sqrt{\alpha_t}\sqrt{\alpha_{t-1}}x_{t-2} + \sqrt{\alpha_t}\sqrt{1 - \alpha_{t-1}}\epsilon_{t-2} + \sqrt{1 - \alpha_t}\epsilon_{t-1}\end{aligned}$$

Let  $G_1 \sim \mathcal{N}(0, \sigma_1^2 I)$ ,  $G_2 \sim \mathcal{N}(0, \sigma_2^2 I)$ , the sum of them gives  $g_2 \sim \mathcal{N}(0, (\sigma_1^2 + \sigma_2^2)I)$ .

# Decreasing training data generation cost

Given a data image  $x_0$ , compute  $x_t$  takes  $t$  sampling on  $q$ . But a simple trick, allows to do only one.

Remember that  $q(x_{t+1} | x_t) := \mathcal{N}(x_{t+1}; \sqrt{1 - \beta_t}x_t, \beta_t I)$ . Let  $\alpha_t = 1 - \beta_t$  and  $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$ .

$$\begin{aligned}x_t &= \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon_{t-1} \\&= \sqrt{\alpha_t}\sqrt{\alpha_{t-1}}x_{t-2} + \sqrt{\alpha_t}\sqrt{1 - \alpha_{t-1}}\epsilon_{t-1} + \sqrt{1 - \alpha_t}\epsilon_{t-1} \\&= \sqrt{\alpha_t\alpha_{t-1}}x_{t-2} + \sqrt{\alpha_t(1 - \alpha_{t-1}) + 1 - \alpha_t}\bar{\epsilon}_t\end{aligned}\tag{1}$$

Let  $G_1 \sim \mathcal{N}(0, \sigma_1^2 I)$ ,  $G_2 \sim \mathcal{N}(0, \sigma_2^2 I)$ , the sum of them gives  $g_2 \sim \mathcal{N}(0, (\sigma_1^2 + \sigma_2^2)I)$ .

# Decreasing training data generation cost

Given a data image  $x_0$ , compute  $x_t$  takes  $t$  sampling on  $q$ . But a simple trick, allows to do only one.

Remember that  $q(x_{t+1} | x_t) := \mathcal{N}(x_{t+1}; \sqrt{1 - \beta_t}x_t, \beta_t I)$ . Let  $\alpha_t = 1 - \beta_t$  and  $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$ .

$$\begin{aligned}x_t &= \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon_{t-1} \\&= \sqrt{\alpha_t}\sqrt{\alpha_{t-1}}x_{t-2} + \sqrt{\alpha_t}\sqrt{1 - \alpha_{t-1}}\epsilon_{t-1} + \sqrt{1 - \alpha_t}\epsilon_{t-1} \\&= \sqrt{\alpha_t\alpha_{t-1}}x_{t-2} + \sqrt{\alpha_t(1 - \alpha_{t-1}) + 1 - \alpha_t}\bar{\epsilon}_t \\&= \sqrt{\alpha_t\alpha_{t-1}}x_{t-2} + \sqrt{1 - \alpha_t\alpha_{t-1}}\bar{\epsilon}_t\end{aligned}\tag{1}$$

Let  $G_1 \sim \mathcal{N}(0, \sigma_1^2 I)$ ,  $G_2 \sim \mathcal{N}(0, \sigma_2^2 I)$ , the sum of them gives  $g_2 \sim \mathcal{N}(0, (\sigma_1^2 + \sigma_2^2)I)$ .

# Decreasing training data generation cost

Given a data image  $x_0$ , compute  $x_t$  takes  $t$  sampling on  $q$ . But a simple trick, allows to do only one.

Remember that  $q(x_{t+1} | x_t) := \mathcal{N}(x_{t+1}; \sqrt{1 - \beta_t}x_t, \beta_t I)$ . Let  $\alpha_t = 1 - \beta_t$  and  $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$ .

We have  $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$ .



# Training

For now, our model is learning  $\mu$  and  $\Sigma$ , i.e. we sample according to

$$p_{\theta}(x_{t-1} \mid x_t) := \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \Sigma_{\theta}(x_t, t))$$

They've found that **fixing  $\Sigma_{\theta}$**  to a constant gives the same result. So,

$$p_{\theta}(x_{t-1} \mid x_t) := \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \sigma_t I)$$

$$p_{\theta}(x_{t-1} \mid x_t) := \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \sigma_t I)$$

The probability for our model to generate  $x_0$  is

$$p_{\theta}(x_0) := \int p_{\theta}(x_{0:T}) dx_{1:T}.$$

# Training

$$p_{\theta}(x_{t-1} \mid x_t) := \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \sigma_t I)$$

The probability for our model to generate  $x_0$  is  $p_{\theta}(x_0) := \int p_{\theta}(x_{0:T}) dx_{1:T}$ . Using negative log likelihood, approximations and computations, we want to minimize:

$$E_q \left[ \frac{1}{2\sigma_t^2} \|\tilde{\mu}_t(x_t, x_0) - \mu_{\theta}(x_t, t)\|^2 \right]$$

where  $\tilde{\mu}$  is the optimal mean that depends on  $x_0$  which we don't know.

# Training

$$p_{\theta}(x_{t-1} \mid x_t) := \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \sigma_t I)$$


The probability for our model to generate  $x_0$  is  $p_{\theta}(x_0) := \int p_{\theta}(x_{0:T}) dx_{1:T}$ . Using negative log likelihood, approximations and computations, we want to minimize:

$$E_q \left[ \frac{1}{2\sigma_t^2} \|\tilde{\mu}_t(x_t, x_0) - \mu_{\theta}(x_t, t)\|^2 \right]$$

where  $\tilde{\mu}$  is the optimal mean that depends on  $x_0$ . Using  $x_t(x_0, \epsilon) = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$  we have a loss we can train on.


# Our results - Gaussian

We have started with Gaussian generation:



# Our results - Gaussian

We have started with Gaussian generation and got satisfying results:



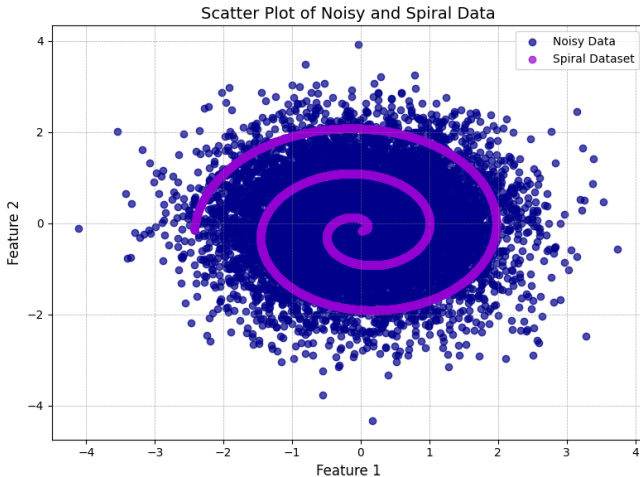
# Our results - Deux gaussiennes

---

A ajouter

# Our results - Spirale

Then we moved to a more complicated dataset, Spirale generation:





# Our results - Spirale

Then we moved to a more complicated dataset, **Spirale generation** and also got satisfying results:

